

Make the first five minutes worthy of the idea.

UX research and product-experience audit of CAOCAP v8.1.1 — onboarding, navigation, visual canvas, project creation, AI assistance, editing, and preview — with prioritized recommendations, an evidence-based issue register, redesigned first-run flows, and a Phase-2 validation plan.

Prepared for: Azzam A. · CAOCAP

Prepared by: Wally · Watermelons Studio

August 2026 · Version 3.0

Contents

- 01 Executive summary
- 02 Product understanding — and the audience question
- 03 Evidence & method
- 04 The first-run funnel, screen by screen
- 05 Findings by product area
- 06 Prioritized issue register
- 07 Recommendations — the redesigned first session
- 08 Live-session annotated screenshot
- 09 Competitive landscape
- 10 Phase-2 research & measurement plan
- 11 Roadmap & success metrics
- 12 Sources

01 Executive summary

CAOCAP's core idea is strong and differentiated: a creative canvas where people learn real software skills by building runnable Mini-Apps with an AI mentor. The building blocks of that vision already exist in v8.1.1 — starter games, daily challenges, an activity streak, a tutorial, and CoCaptain itself. What blocks the product today is sequencing: the first session spends its opening minutes on narrative, interface concepts, and a professional review workflow, and delays the one thing that would make a new user stay — a visible change to a running app, made by them.

The first session in one sentence: users are asked to review a pull request before they are allowed to play.

Three numbers frame the stakes:

- 3.5 / 5 App Store rating from only 13 ratings since 2019 — interest exists at the top of the funnel, but the first session is not converting it.
- Industry benchmarks: roughly 77% of users abandon an app within 3 days of install, and only ~8.4% of users complete onboarding flows at all (AppsFlyer via VWO 2025; scandiweb 2026). The first session is the highest-leverage surface in the product.
- ~60 seconds is the window in which a new user should feel core value (lowcode.agency 2026). CAOCAP's current first visible outcome — a renamed text string, gated behind a code diff — arrives minutes late and lands flat.

The five moves that matter most:

- P0 — Rebuild the first session around one activation moment: the user makes a visible change to a running Mini-App within 2–3 minutes. Every other screen serves that moment or moves after it.
- P0 — Unlock CoCaptain for Mission #1. The AI mentor is the product's differentiator; it is currently presented as paywalled before a new user experiences any value. Monetize depth, never the first win.
- P0 — Remove the git-style diff from the beginner path. Changes apply instantly to a live preview with a prominent Undo; reading diffs becomes a deliberate advanced lesson.
- P0 — Resolve the audience question. Independent reviewers read CAOCAP as three different products (a children's app, a general learning app, a builder for creative adults). Until one primary audience is declared, every copy, pricing, and UI decision is unanchored.
- P1 — Make the canvas discoverable and calm: pressed states and affordances on every node, no camera movement without user intent, fit-to-view, and a beginner home of 3–4 nodes instead of ten.

02 Product understanding — and the audience question

2.1 What CAOCAP is

CAOCAP is a mobile-first creative environment in which projects live as nodes on an infinite spatial canvas. Users build and modify small web apps (HTML/CSS/JS) — games, tools, experiments — assisted by a selectable AI co-pilot (CoCaptain / CoStar) that chats, executes build requests, and explains decisions. A command palette (⌘) acts as command launcher and AI prompt field. Ready Mini-Apps (Pac-Man, XO, Hello World) demonstrate what can be built; a seven-step tutorial walks new users from the canvas into a first AI-assisted edit; daily building challenges and an activity heatmap supply habit mechanics. Free tier with Pro at \$9.99/month or \$79.99/year. Requires iOS 26+. English and Arabic.

2.2 The audience question — a strategic finding

The product's own vision document targets “creative builders”: designers, founders, students, and product thinkers. The onboarding, however, opens with story-world narrative and a mascot-like co-pilot choice, while the working surface exposes professional developer conventions — nodes, sub-canvases, code diffs, an ‘Agent’ toggle, a ‘Gemini cloud’ model selector, a Mac ⌘ glyph on a phone.

The result is measurable in how professionals read the product: of the independent expert reviews commissioned for this project, one audited CAOCAP as a children's learning product, one as a general consumer learning app, and this report's evidence base (the vision document) points at creative adults. Three auditors, three products.

If three professional reviewers cannot agree on who CAOCAP is for, a new user deciding in 30 seconds has no chance. Declaring one primary audience is a PO product decision that precedes visual polish.

Our recommendation, argued throughout this report: commit to creative beginners (teens and adults) as the primary audience, with a distinct kids mode as a later, deliberate product line if desired. The tone can stay playful; the workflow must stop assuming professional literacy.

03 Evidence & method

Findings in this report are triangulated across four independent evidence streams; no finding rests on a single source.

Stream	What it covers	Key artifacts
A · Store & history forensics	7 years of releases, positioning shifts, user reviews, developer responses	App Store listing v1.1→v8.1.1; all public reviews; IAP structure
B · Live device session	Hands-on session on v8.1.1 on the client-provided install (Aug 3, 2026)	Annotated home-canvas capture (Section 08)
C · Fresh-install walkthroughs	Complete 14-screen first-run funnel on clean installs, corroborated by two independent expert walkthroughs shared with the project	Screen-by-screen funnel (Section 04)
D · External benchmarks	Onboarding and retention norms for mobile learning products	AppsFlyer/VWO 2025; scandiweb 2026; Adaptly A/B data; Duolingo-pattern references

Heuristic frame: Nielsen's 10 usability heuristics, applied across the seven areas in the client brief (onboarding, navigation, visual canvas, project creation, AI assistance, editing, preview). Severity: CRITICAL blocks the core loop or causes abandonment · MAJOR significant recurring friction · MINOR polish. Limitations: the live session covered the home canvas but not deep editing flows (unstable desktop mirroring environment); the fresh-install funnel is corroborated by independent walkthroughs rather than captured first-hand. Phase 2 (Section 10) closes both gaps with moderated tests and full instrumentation.

04 The first-run funnel, screen by screen

A brand-new user's first session, read as a funnel (~3 minutes, English UI). The pattern: narrative and interface concepts are front-loaded; the value moment is back-loaded and then under-delivered.

#	Screen	What it asks of the user	Assessment
1-3	Story slides ("Countdown to launch", "One small step"...)	Absorb brand narrative; no input collected	3 passive screens; adult, metaphor-heavy copy; nothing adapts
4	Choose your co-pilot (CoCaptain / CoStar)	Pick a character with no stated capability difference	Choice feels arbitrary; answer changes nothing visible
5	"What's your coding level?" slider with HTML/CSS/JS logos	Self-assess against file-format logos on a continuous slider	Wrong control, jargon iconography; answer routes nothing downstream
6	"Tap to launch your coding journey"	Transition tap	Pure ceremony before any value
7	Main canvas — ten nodes incl. Pro Subscription, Settings	Orient among 10 mixed-purpose live tap targets	Cognitive dump; the real next step (Tutorial) competes with 9 distractors; Pro card visible pre-value
8-10	Tutorial tooltips ×3 for the ⌘ command palette	Learn a dual-purpose input via reading	An input that needs 3 tooltips is the problem, not the user; ⌘ is a desktop glyph on a phone
11	CoCaptain chat: 6-line greeting; edit-task prompt	Read a text wall; system jargon visible ('Agent', 'Gemini cloud')	Mentor moment becomes a lecture; implementation details leak
12-13	Pending changes: git-style diff, Apply / Reject / View	Review red/green code blocks to approve a text rename	First creative act gated behind a professional review ritual; grammar bug ("1 changes need your review")
14	Celebration → "turn it into a mini game from Help"	Feel rewarded for a renamed string; continue via documentation	Reward doesn't match a creation; the next step exits to docs — sessions end here

Funnel verdict: the session teaches the interface before it lets the user create, asks personalization questions whose answers change nothing, and spends its single celebration on a renamed heading. Every cliff in this table is addressed by a specific change in Section 07.

05 Findings by product area

[CRITICAL] F1 — Onboarding — value is never demonstrated, then a diff blocks the win

- Three story slides deliver mood, not meaning; the product's actual promise (build real apps with an AI mentor) is never stated in benefit language.
- The only personalization asked (co-pilot, skill slider) visibly changes nothing afterwards — a documented missed lever: published A/B data shows personalization that shapes the flow lifts trial starts ~8.5% and paid conversion ~17% (Adapty).
- The first visible outcome is a renamed text string, reached through a git-style diff with Apply/Reject. Effort goes into bureaucracy, not creation — the exact inversion of the endowed-progress effect that makes playful learning apps retain.
- Signup and Pro messaging appear before any value is felt; the celebration screen exits into documentation (Help).

[CRITICAL] F2 — Visual canvas — undiscoverable, crowded, and camera-unstable

- First landing is a ten-node canvas mixing creations with Profile, Settings, Pro Subscription, App Icon — a file browser presented as a first impression. No product's aha moment is its file browser.
- No affordance signals what is tappable, draggable, or zoomable; in our live session, card taps produced no pressed state and no visible response.
- Documented history of the camera acting without user intent (auto-zoom on add, fly-to stutter — store review + release notes). Rule: the camera belongs to the user; add fit-to-view.
- Requirements, structure, styling, logic, and preview share one infinite surface without progressive disclosure — novices see everything and understand nothing.

[CRITICAL] F3 — AI assistance (CoCaptain) — the differentiator is gated, verbose, and leaky

- The home canvas presents CoCaptain as locked behind Pro (“Unlock CoCaptain & Premium Features”) before a new user experiences any value — paywalling the reason the product exists.
- The first AI moment is a six-line greeting plus visible system internals (‘Agent’ toggle, ‘Gemini cloud’ model selector). A mentor's first line should be one sentence and a suggested action.
- Release notes document silent AI response failures (v7.1.1 “assistant responses could occasionally fail to parse”); the assistant needs explicit loading/streaming/failure states with retry.
- AI edits should always land as a reviewable card with ‘Apply’ + ‘Explain why’ in beginner-readable language — this is CAOCAP's honest differentiator against prompt-to-app tools, and it must be visible in the UI, not aspirational.

[MAJOR] F4 — Navigation & IA — three pivots left mixed metaphors and no orientation

- Nodes, sub-canvases, missions, SRS notes, code editors, and previews coexist without one clear hierarchy — residue of three product re-architectures (MindMap tool → agent IDE → learning canvas).
- No persistent answer to: where am I, what am I working toward, how do I get back? A mission header with progress and a breadcrumb solves this cheaply.
- The ☿ command palette is a desktop convention carrying two mental models (commands + AI prompt) behind one unlabeled glyph; the tutorial needs three tooltips to explain one send. In beginner mode, the chat with the co-pilot should simply be the interface.
- Features appear and disappear across versions (daily challenges, badges, publishing) — learnability and trust erode each time.

[MAJOR] F5 — Project creation — cold start against the product's own philosophy

- Creation defaults toward blank/abstract structure while the product's stated principle is learning inside real apps. Runnable starters exist (Pac-Man, XO — verified in our live session) but are not the mandatory front door.
- “What do you want to build?” — the single most valuable personalization question for a builder product — is never asked. The chosen answer should select the first template and make the session feel personal.

- Account creation sits before value. Move signup after the first win, framed as saving your creation — the pattern publicly credited with ~20% DAU lift when Duolingo moved signup behind the first lesson.

[MAJOR] F6 — Editing & trust — hidden surfaces, invisible saves

- Store reviews document users unable to find where to edit content or styling at all; every node needs a one-tap contextual inspector, and every gesture a visible equivalent control.
- v7.1.1 shipped “prevent data loss” fixes — users have lost work. Autosave must be visible (“Saved ✓”), undo/redo prominent, and version snapshots reachable.
- Grammar and copy seams on flagship surfaces (“1 changes need your review”) undermine perceived quality at the exact moment of highest attention.

[MAJOR] F7 — Preview — the magic loop is not persistent

- Edit → see result is the product's core magic, yet preview is a destination rather than a companion. It should be one tap away everywhere, split-view on iPad, and update live on every change.
- Run/inspect states need one obvious control; device frames matter for users building web Mini-Apps.

[POSITIVE] F8 — Foundations already in place (build on these)

- Runnable starter Mini-Apps (Pac-Man, XO) labeled mobile-ready — the never-blank-canvas foundation exists.
- Daily building challenges (0/3) and an activity heatmap — habit-loop mechanics are half-built and only need connecting to a mission funnel.
- A guided tutorial exists as a sub-canvas — the right concept, currently optional; Section 07 makes it the spine of the first session.
- Direct founder support (“Message Azzam directly” via WhatsApp) — charming at this stage; move into Help to declutter the home canvas.

06 Prioritized issue register

ID	Area	Issue	Evidence	Sev.	Prio
U1	Onboarding	No value proposition stated in first-run; 3 passive story slides	Fresh-install funnel	Crit	P0
U2	Onboarding	First creative act gated behind git-style diff (Apply/Reject)	Fresh-install funnel	Crit	P0
U3	Onboarding	Personalization (co-pilot, level slider) changes nothing downstream	Fresh-install funnel	Maj	P1
U4	Onboarding	"What do you want to build?" never asked	Flow-level gap	Maj	P1
U5	Onboarding	Celebration rewards a renamed string; next step exits to Help docs	Fresh-install funnel	Maj	P1
U6	Monetization	CoCaptain presented as Pro-locked before first value	Live session	Crit	P0
U7	Monetization	Signup before value; paywall timing premature	Funnel + store IAP	Maj	P1
U8	Canvas	First landing = 10 mixed-purpose nodes incl. Settings & Pro	Live session + funnel	Crit	P0
U9	Canvas	No affordances / pressed states; taps gave no visible response	Live session	Crit	P0
U10	Canvas	Camera moves without user intent; zoom bugs; no fit-to-view	Reviews + release notes	Crit	P0
U11	Navigation	☞ palette: desktop glyph, dual mental model, needs 3 tooltips	Fresh-install funnel	Crit	P0
U12	Navigation	No mission header / breadcrumb / orientation cues	Heuristic + live session	Maj	P1
U13	AI	6-line greeting text wall; 'Agent'/'Gemini cloud' jargon exposed	Fresh-install funnel	Maj	P1
U14	AI	History of silent AI failures; no visible states/retry	Release notes v7.1.1	Maj	P1
U15	AI	AI edits lack 'Apply + Explain why' framing in beginner language	Vision-doc gap	Maj	P1
U16	Creation	Runnable starters exist but are not the mandatory front door	Live session	Maj	P1
U17	Editing	Edit surfaces undiscoverable (content/styling)	Store reviews	Crit	P0
U18	Editing	Data-loss history; autosave not visible; undo not prominent	Release notes v7.1.1	Maj	P1
U19	Preview	Preview not persistent; entry unclear	History + heuristic	Maj	P1
U20	Strategy	Audience undeclared — read as 3 different products by 3 reviewers	Cross-report analysis	Crit	P0
U21	Store	Placeholder artwork; conflicting subtitle; 3.5★; iOS 26+ only	Store listing	Maj	P1
U22	Quality	Copy/grammar seams on flagship screens ("1 changes need...")	Fresh-install funnel	Min	P2
U23	A11y	No accessibility declaration; canvas VoiceOver behavior unknown	Store listing	Min	P2
U24	Home	System cards (Profile/Settings/Sub/App Icon) mixed with content	Live session	Maj	P1

P0 = fix before any marketing push · P1 = next release cycle · P2 = polish backlog. Nine issues are P0; all nine are concentrated in the first session and the canvas — one coherent redesign addresses them together.

07 Recommendations — the redesigned first session

One organizing rule: the user makes a visible change to a running Mini-App within the first 2-3 minutes. Every screen either serves that moment or moves after it.

7.1 The mental-model shift

First minute — today	First minute — redesigned
"Beautiful space theme... but what is this app, and what am I supposed to do?"	"I picked a game, it's already running, and I just changed it myself. The AI explained why it worked."
Reads narrative → picks a mascot → self-assesses on a slider → lands among 10 nodes → reads 3 tooltips → reviews a diff	Sees one promise line → picks what to build → game opens running → makes 3 tiny guided changes → celebrates → saves to continue

7.2 Seven changes, in order of leverage

C1 — Collapse the opening to one promise + one trust line.

One screen: the co-pilot introduces itself in a single sentence and states the outcome ("You'll have a real mini-game running in 3 minutes — no experience needed"). All further narrative becomes short reactions to the user's own choices.

C2 — Ask only questions that change the product.

Three one-tap screens with a progress bar: What do you want to build? (picture cards — selects the first template) · Have you built before? (plain-language levels — routes hand-holding) · Daily building goal (converts into a payoff projection). Cut the co-pilot beauty contest or give the choice stated capability differences.

C3 — Land in a running Mini-App, never on the canvas.

The chosen template opens already running, full screen. The canvas is introduced after the first win as 'home', and in beginner mode it contains 3-4 nodes: your creations, Daily, Tutorial. System items (Profile, Settings, Subscription, App Icon) move to a standard settings sheet.

C4 — Make the first build a guided mission of instant wins.

3-4 micro-tasks in a lesson rhythm (change a color, make it jump higher, rename it, add a sound), each applying instantly to the live game with a prominent Undo. No diff in beginner mode — reading and approving code becomes a deliberate advanced lesson later. Close with a stats screen (changes made, time, streak spark) and the game playable full screen.

C5 — Free Mission #1 with CoCaptain; signup and Pro after the win.

The mentor is free for the first mission for everyone; monetize depth (advanced missions, unlimited AI, publishing). Signup is framed as saving your creation. Pro messaging never blocks the first session.

C6 — Give CoCaptain interaction discipline.

First line = one sentence + one suggested action chip. Visible thinking/streaming/failure states with retry. Every AI edit lands as a card: what changed, in plain language, with Apply and Explain-why. Hide 'Agent', model names, and internals behind an advanced toggle.


C7 — Fix the seams that price the product down.

Replace the ☹️ glyph with the co-pilot's face as the single entry point; pressed states on every tappable card; camera never moves uninvited + fit-to-view; visible "Saved ✓"; grammar pass on flagship screens; store listing refresh (real screenshots of the mission loop, one consistent subtitle).

7.3 The redesigned flows


First-Run Experience - Current vs. Improved

CURRENT (drop-off at every step)



Evidence: "no manual... unintuitive" / "can't figure out how to use it - and I'm a Software Engineer" (App Store reviews)


IMPROVED (first win in under 5 minutes)



Principles: value before signup / learning IS the onboarding / one mission replaces the manual / celebrate the first visible result

KPI funnel to instrument: install > demo started > mission 1 completed > signup > D1/D7 return


Improved Core Loop - Guided Building Missions



Design rules: mission <= 5 min / never a blank canvas / AI edits always a reviewable diff / explanation arrives when needed, not before

7.4 The screen that carries the loop


Improved Canvas Layout - Wireframe (iPad / iPhone landscape)



Key fixes: every object is tappable with a visible edit surface / camera never moves without user intent / mission, canvas, result on one screen

08 Live-session annotated screenshot

Captured during our hands-on session on v8.1.1 (client-provided install, Aug 3, 2026). Numbered annotations map to register IDs U6, U8, U9, U16, U24 and the positives in F8.



CAOCAP v8.1.1 — Home Canvas

Live session · Aug 3, 2026 · via iPhone Mirroring

- 1 CoCaptain — the core AI mentor — is locked behind a Pro subscription: paywalled before the user gets any first value.
- 2 System cards (Profile, Settings, Subscription, App Icon) are mixed with learning content on the same canvas — no hierarchy.
- 3 Tutorial exists but is optional ('Tap to open'). It should BE the forced guided first-run, not a card you may never tap.
- 4 GOOD — Real runnable starter Mini-Apps exist (Pac-Man, XO, labeled 'mobile-ready'). Right foundation for never-blank-canvas.
- 5 GOOD — Daily building challenges (0/3) and an Activity heatmap: the habit-loop foundations are already built.
- 6 'WhatsApp — Message Azzam directly': charming high-touch support for early stage, but clutters the home canvas. Move into Help.
- 7 Unlabeled floating command button — a mystery-meat primary control. Needs a label or progressive tooltip.
- 8 No mission header or orientation cues. In our session, card taps produced no pressed-state and no visible response.

Caveat: tested through iPhone Mirroring on a memory-constrained Mac; tap unresponsiveness (8) needs on-device confirmation. Absence of a pressed/feedback state is a defect regardless of root cause.

09 Competitive landscape

Product	Model	What they nail	What CAOCAP can own
Mimo	5-min gamified lessons + AI, career paths	Habit loop, tiny wins, polished onboarding	Real spatial building vs. linear lessons
Sololearn	Courses + community + code battles	Social proof, generous free tier	A mentor that sees your whole project
Swift Playgrounds	Apple's puzzle-based Swift learning	Native polish, zero-setup magic	Web tech + AI mentorship + open canvas
Brilliant	Interactive concept learning	Explaining ideas through interaction	Concepts taught inside a real product you keep
Duolingo (pattern reference)	Mascot-led micro-lesson habit machine	Every onboarding screen collects input or delivers product; signup after first lesson	Same discipline applied to building, not translating
Prompt-to-app tools (Bolt/v0-class)	Generate whole apps from prompts	Instant results	Authorship & understanding — the anti-dependence position

Strategic position: visible, learnable software. Diff-cards-with-explanations, just-in-time concepts, and missions are the interface expression of that position — they are what prompt-to-app tools structurally cannot copy.

10 Phase-2 research & measurement plan

10.1 Research questions

- RQ1 — Can a first-time creative beginner reach a visible win unaided, and where exactly do they stall?
- RQ2 — Do users understand the canvas metaphor (nodes, sub-canvases, layers)? What mental model do they form?
- RQ3 — Is CoCaptain experienced as a mentor or a vending machine? Do users trust and learn from its edits?
- RQ4 — What triggers drop-off in session 1 and week 1?

10.2 Method

- 5–8 moderated sessions (60 min) with creative beginners — designers, founders, students with an app idea; include 2 Arabic-first users. Tasks: free exploration → create → complete Mission 1 → change text + color → add behavior via CoCaptain → preview and narrate.
- Metrics: task success, time-to-first-win, assists, SUS, 3 emotion words. Sessions recorded; findings mapped back to the register IDs in Section 06.

10.3 Instrument the funnel before redesigning

Event	Definition
onboarding_started	First screen rendered on a fresh install
build_intent_chosen	User selects what to build (new C2 screen)
template_opened_running	First Mini-App visible and running
first_change_applied	PRIMARY KPI — user-initiated change visible in the running app
mission1_completed	All micro-tasks of the first mission done
signup_completed	Account created (should sit after mission1_completed)
d1_return / d7_return	User returns day 1 / day 7

Hypothesis to confirm with data: today's two largest cliffs are the ten-node canvas landing and the diff review. Instrumentation settles this before further design investment.

11 Roadmap & success metrics

Horizon	Focus	Ship
0–30 days	Stop the bleeding (P0)	Redesigned first session (C1–C4) · CoCaptain free for Mission 1 · diff removed from beginner path · pressed states + camera rules + fit-to-view · beginner home of 3–4 nodes · audience decision made and written down
30–60 days	Trust & loop (P1)	CoCaptain states + explain-cards (C6) · signup-after-win + streak ladder · contextual inspector on every node · persistent preview · mission header · store listing overhaul
60–90 days	Delight & depth (P2)	Motion/sound celebration system · advanced mode (diff lessons, model controls) · accessibility pass · Arabic-first QA · publishing/share loop

North-star: % of new installs reaching `first_change_applied` in the first session. **Guardrails:** D1/D7 return, missions per week, Pro conversion after Mission 3.

12 Sources

- App Store listing, reviews, version history v1.1→v8.1.1 — apps.apple.com/us/app/caocap/id1447742145
- Live device session on v8.1.1, client-provided install, Aug 3 2026 (annotated capture, Section 08)
- Fresh-install first-run walkthroughs corroborated by two independent expert reviews shared with the project (Aug 2026)
- CAOCAP product vision — azzam.ai/caocap/learn (“Azzamification”: audience, core loop, principles); site: azzam.ai/caocap; open source: github.com/CAOCAP
- Benchmarks: AppsFlyer retention data via VWO (2025); scandiweb onboarding completion (2026); lowcode.agency 60-second value guidance (2026); Adapty personalization A/B results; Duolingo onboarding pattern documentation (First Round Review)
- Competitive references: Mimo, Sololearn, Swift Playgrounds, Brilliant public materials and 2026 comparisons (coursefacts.com, hackr.io)